

# Chapter 4

## Network Layer:

### The Data Plane


A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016  
J.F. Kurose and K.W. Ross, All Rights Reserved



*Computer  
Networking: A Top  
Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Network Layer: Data Plane 4-1

## Chapter 4: outline

### 4.1 Overview of Network layer

- data plane
- control plane

### 4.2 What's inside a router

### 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing

### 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

Network Layer: Data Plane 4-2

# Chapter 4: network layer

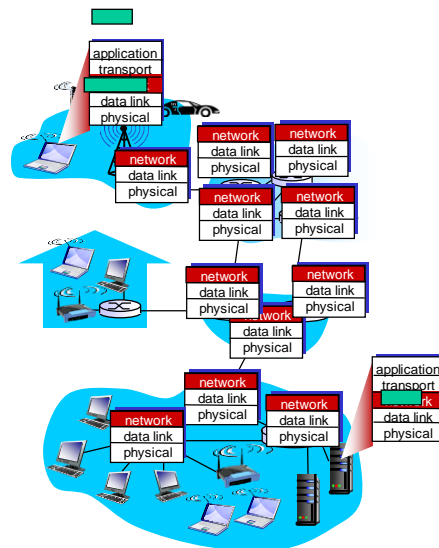
## *chapter goals:*

- understand principles behind network layer services, focusing on data plane:
  - network layer service models
  - forwarding versus routing
  - how a router works
  - generalized forwarding
- instantiation, implementation in the Internet

Network Layer: Data Plane 4-3

## Network layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it



Network Layer: Data Plane 4-4

## Two key network-layer functions

### *network-layer functions:*

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

### *analogy: taking a trip*


- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination

Network Layer: Data Plane 4-5

## Network layer: data plane, control plane

### *Data plane*

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
- forwarding function




### *Control plane*

- network-wide logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers

Network Layer: Data Plane 4-6

## Per-router control plane

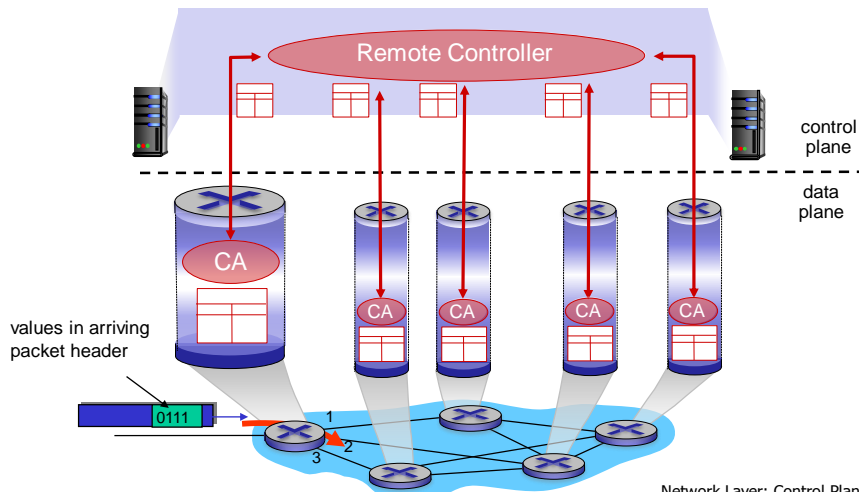
Individual routing algorithm components *in each and every router* interact in the control plane



Network Layer: Control Plane 5-7

## Logically centralized control plane

A distinct (typically remote) controller interacts with local control agents (CAs)



Network Layer: Control Plane 5-8

## Chapter 4: outline

### 4.1 Overview of Network layer

- data plane
- control plane

### 4.2 What's inside a router

### 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing


### 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

Network Layer: Data Plane 4-9

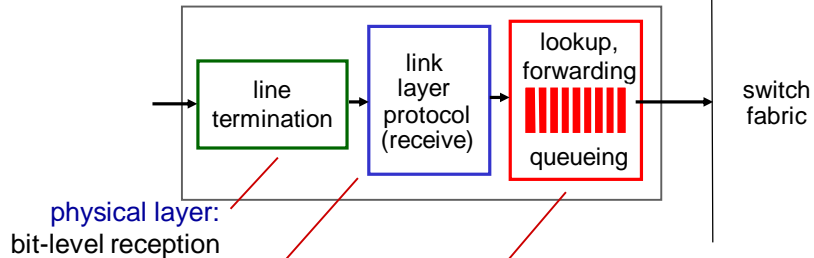
## Router architecture overview

- high-level view of generic router architecture:



Network Layer: Data Plane 4-10

## Input port functions



physical layer:  
bit-level reception


data link layer:  
e.g., Ethernet  
see chapter 6

### decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory ("*match plus action*")
- goal: complete input port processing at 'line speed'
- queuing: if datagrams arrive faster than forwarding rate into switch fabric

Network Layer: Data Plane 4-11

## Input port functions



physical layer:  
bit-level reception

data link layer:  
e.g., Ethernet  
see chapter 6

### decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory ("*match plus action*")
- *destination-based forwarding*: forward based only on destination IP address (traditional)
- *generalized forwarding*: forward based on any set of header field values

Network Layer: Data Plane 4-12

## Destination-based forwarding

*forwarding table*

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Q: but what happens if ranges don't divide up so nicely?

Network Layer: Data Plane 4-13

## Longest prefix matching

*longest prefix matching*

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

DA: 11001000 00010111 00010110 10100001

which interface?

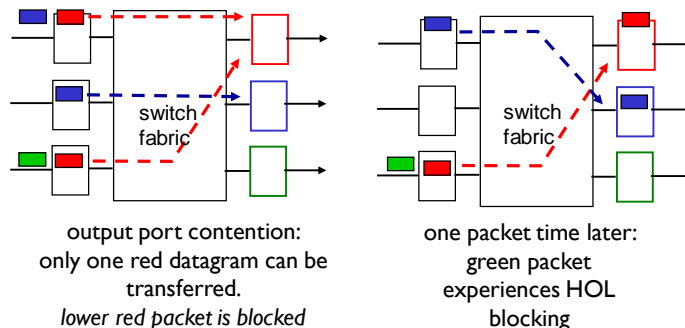
DA: 11001000 00010111 00011000 10101010

which interface?

Network Layer: Data Plane 4-14

## Input port queuing


- fabric slower than input ports combined -> queueing may occur at input queues
  - *queueing delay and loss due to input buffer overflow!*
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward



Network Layer: Data Plane 4-15

## Output ports

*This slide is HUGE important!*



- **buffering** required from fabric faster rate
 


Datagram (packets) can be lost due to congestion, lack of buffers
- **scheduling** datagrams
 

Priority scheduling – who gets best performance, network neutrality

Network Layer: Data Plane 4-16



## Output port queueing



- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

Network Layer: Data Plane 4-17

## Chapter 4: outline

### 4.1 Overview of Network layer

- data plane
- control plane

### 4.2 What's inside a router

#### 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing

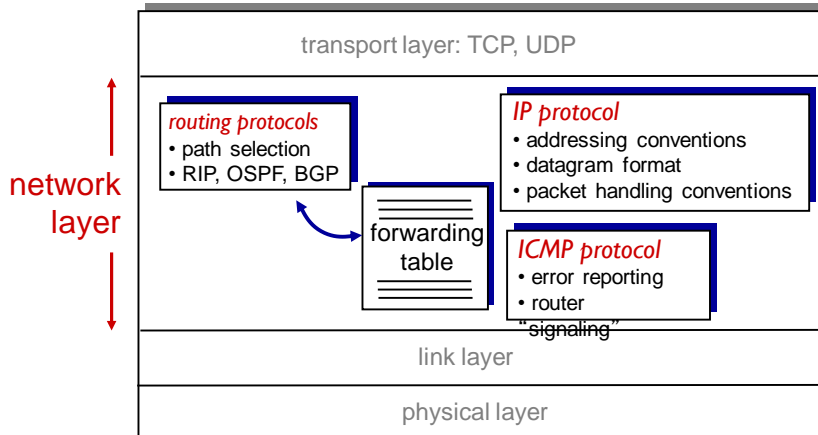
### 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

Network Layer: Data Plane 4-18

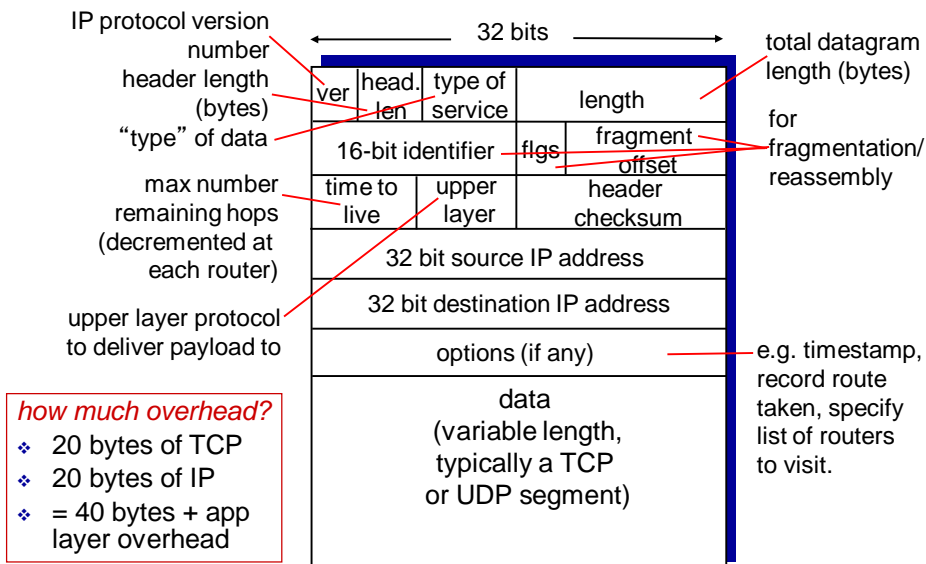
# The Internet network layer

host, router network layer functions:



Network Layer: Data Plane 4-19


## IP datagram format



Network Layer: Data Plane 4-20

## IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided ("fragmented") within net
  - one datagram becomes several datagrams
  - "reassembled" only at final destination
  - IP header bits used to identify, order related fragments



Network Layer: Data Plane 4-21

## IP fragmentation, reassembly

*example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

1480 bytes in  
data field

offset =  
1480/8

length	ID	fragflag	offset
=4000	=x	=0	=0

*one large datagram becomes  
several smaller datagrams*

length	ID	fragflag	offset
=1500	=x	=1	=0

length	ID	fragflag	offset
=1500	=x	=1	=185

length	ID	fragflag	offset
=1040	=x	=0	=370

Network Layer: Data Plane 4-22

## Chapter 4: outline

### 4.1 Overview of Network layer

- data plane
- control plane

### 4.2 What's inside a router

### 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing


### 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

Network Layer: Data Plane 4-23

## IP addressing: introduction

- **IP address:** 32-bit identifier for host, router interface
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP addresses associated with each interface**



Network Layer: Data Plane 4-24

## IP addressing: introduction


**Q:** how are interfaces actually connected?

**A:** we'll learn about that in chapter 6, 7.

**A:** wired Ethernet interfaces connected by Ethernet switches

**For now:** don't need to worry about how one interface is connected to another (with no intervening router)

**A:** wireless WiFi interfaces connected by WiFi base station



Network Layer: Data Plane 4-25


## Subnets

### ■ IP address:

- subnet part - high order bits
- host part - low order bits

### ■ what's a subnet?

- device interfaces with same subnet part of IP address
- can physically reach each other *without intervening router*




network consisting of 3 subnets

Network Layer: Data Plane 4-26

# Subnets

## recipe

- to determine the subnets, detach each interface from its host or router, creating islands of isolated networks
- each isolated network is called a **subnet**




subnet mask: /24

Network Layer: Data Plane 4-27

# Subnets

## how many?




Network Layer: Data Plane 4-28

## IP addressing: CIDR

### CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



Network Layer: Data Plane 4-29

## IP addresses: how to get one?

**Q:** How does a *host* get IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
  - “plug-and-play”

Network Layer: Data Plane 4-30

## IP addresses: how to get one?

**Q:** how does *network* get subnet part of IP addr?

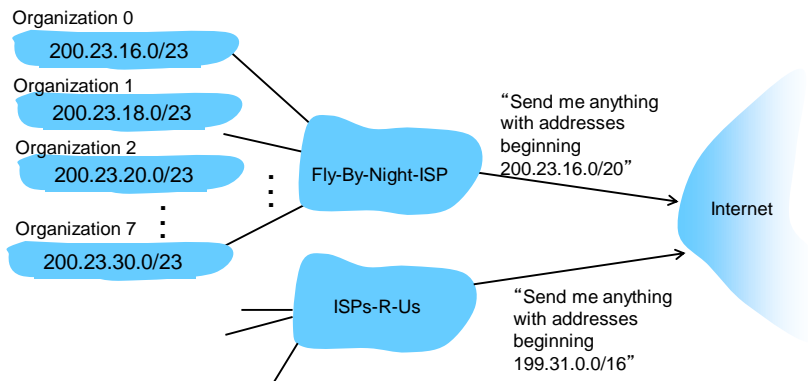
**A:** gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>00010010</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>00010100</u>	00000000	200.23.20.0/23
...	....	....	....	....	....
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>00011110</u>	00000000	200.23.30.0/23

Network Layer: Data Plane 4-31

## Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:

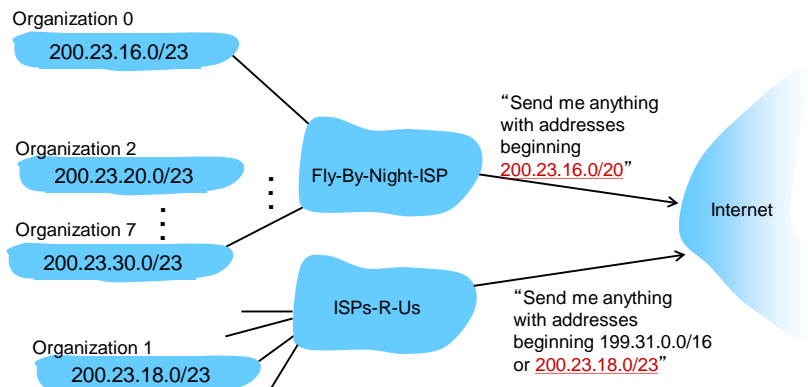


Network Layer: Data Plane 4-32



## Hierarchical addressing: more specific routes

ISPs-R-Us has a more specific route to Organization 1



Network Layer: Data Plane 4-33

## IP addressing: the last word...

**Q:** how does an ISP get block of addresses?

**A: ICANN:** Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes

Network Layer: Data Plane 4-34

## Chapter 4: outline

### 4.1 Overview of Network layer

- data plane
- control plane

### 4.2 What's inside a router

### 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing

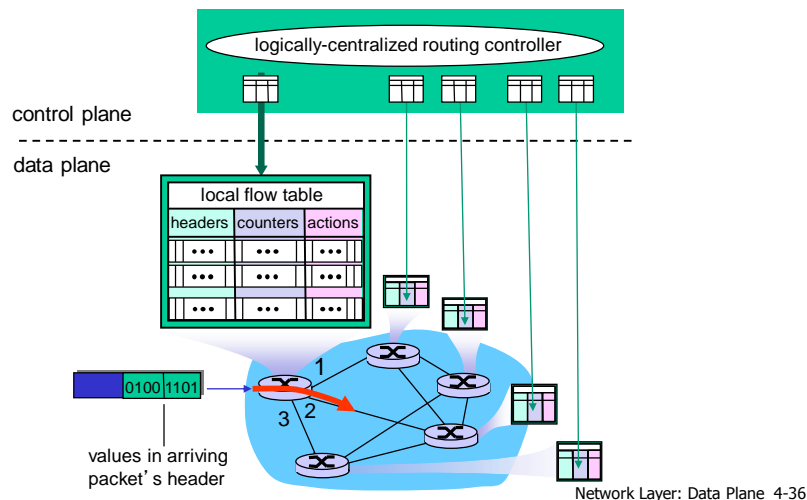
### 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

Network Layer: Data Plane 4-35


## Generalized Forwarding and SDN

Each router contains a *flow table* that is computed and distributed by a *logically centralized routing controller*



## OpenFlow data plane abstraction

- *flow*: defined by header fields
- generalized forwarding: simple packet-handling rules
  - *Pattern*: match values in packet header fields
  - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
  - *Priority*: disambiguate overlapping patterns
  - *Counters*: #bytes and #packets




*Flow table in a router (computed and distributed by controller) define router's match+action rules*

Network Layer: Data Plane 4-37

## OpenFlow data plane abstraction

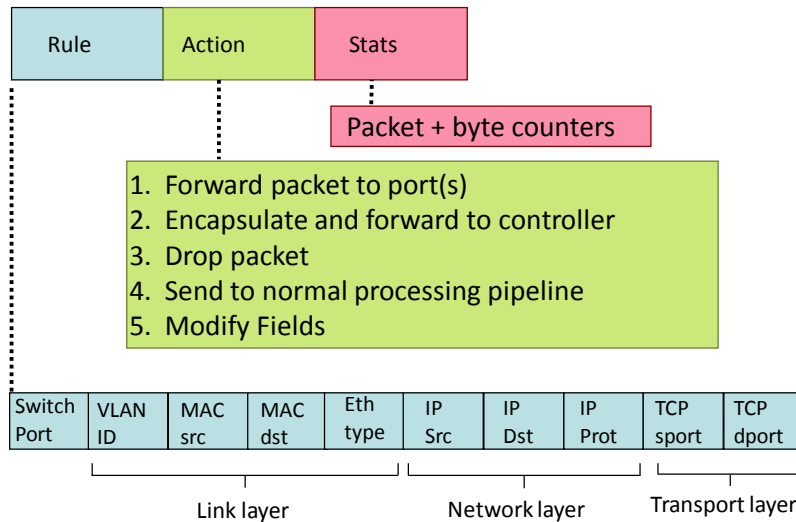
- *flow*: defined by header fields
- generalized forwarding: simple packet-handling rules
  - *Pattern*: match values in packet header fields
  - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
  - *Priority*: disambiguate overlapping patterns
  - *Counters*: #bytes and #packets



\* : wildcard

1. src=1.2.\*.\*, dest=3.4.5.\* → drop
2. src = \*.\*.\*, dest=3.4.\*.\* → forward(2)
3. src=10.1.2.3, dest= \*.\*.\*.\* → send to controller

# OpenFlow: Flow Table Entries



## Examples

### Destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	51.6.0.8	*	*	*	port6

*IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6*

### Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	*	*	*	*	22	drop

*do not forward (block) all datagrams destined to TCP port 22*

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	128.119.1.1	*	*	*	*	drop

*do not forward (block) all datagrams sent by host 128.119.1.1*

# Examples

## Destination-based layer 2 (switch) forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	22:A7:23:11:E1:02	*	*	*	*	*	*	*	*	port3

*layer 2 frames from MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3*

Network Layer: Data Plane 4-41


# OpenFlow abstraction

- **match+action**: unifies different kinds of devices
- Router
  - **match**: longest destination IP prefix
  - **action**: forward out a link
- Switch
  - **match**: destination MAC address
  - **action**: forward or flood
- Firewall
  - **match**: IP addresses and TCP/UDP port numbers
  - **action**: permit or deny
- NAT
  - **match**: IP address and port
  - **action**: rewrite address and port

Network Layer: Data Plane 4-42

## OpenFlow example

*Example:* datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2



## Chapter 4: done!

4.1 Overview of Network layer: data plane and control plane

4.2 What's inside a router

4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing

4.4 Generalized Forward and SDN

- match plus action
- OpenFlow example

*Question:* how do forwarding tables (destination-based forwarding) or flow tables (generalized forwarding) computed?

*Answer:* by the control plane (next chapter)